

CS 6068
Parallel Computing
Fall 2014
Lecture 3 – Sept 15

Prof. Fred Annexstein @proffreda
fred.annexstein@uc.edu
Office Hours: 11-12 MW or by appointment
Tel: 513-556-1807
Meeting: Mondays 6:00-8:50PM Baldwin

Week 3 : Plan

- Review – Basic Parallel Operations

Scatter, Gather, Reduce, Prefix-Scan, Histogram

- More Common Communication Patterns

Compact/Filter

- **Segmented Scan with Apps**

Sparse Matrix Products

CSR format

- Parallel Sorting

- Sorting Networks – **Bitonic Sort**

- **Homework #4 Sorting Arrays for red-eye removal**

Compact / Filter

- returns a sequence consisting of those items from the sequence for which *predicate(item)* is true. If *sequence* is a string or tuple, the result will be of the same type; otherwise, it is always a list. For example, to compute some small primes:
- ```
>>> def prd(x): return x % 2 != 0 and x % 3 != 0
```
- ```
>>> compact(prd, range(2, 20))
```
- ```
[5, 7, 11, 13, 17, 19]
```

# Relation of Scatter to Compact

- We assume that the filter predicate can be applied to each element in the list in parallel constant (fast) time, resulting in a bit vector of Boolean values.
- |   |   |   |   |   |   |   |   |    |    |    |    |    |    |    |    |    |    |   |
|---|---|---|---|---|---|---|---|----|----|----|----|----|----|----|----|----|----|---|
| 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 |   |
| 0 | 0 | 0 | 1 | 0 | 1 | 0 | 0 | 0  | 1  | 0  | 1  | 0  | 0  | 0  | 0  | 1  | 0  | 1 |
- Let's look at the running sum of these Boolean values.... 000112222334444556

# Scatter Operations

- Recall that Scatter is an data communication operation that moves or permutes data based on an array of location addresses
- The index of each element in the filtered or compacted list is the running sum, or more precisely, the exclusive sum scan.
- For example, we see that 17 is in the 4th list position since its running sum value is 5 which is 1 greater than its left neighbor which is 4.

# Quiz

- Complexity: What is the Work and Step Complexity of the Compact Operation on arrays?
- Let's consider Irregular Workloads....
- Can you generalize Compact so that elements can allocate space on a per filtered element basis?  
For example, suppose each thread needs to dynamically allocate 0-10 memory locations

# Review Complexity Analysis of Scan

## Recursive Version

```
def prefix(add,x):
 if len(x)== 1: return x
 else:
 firsthalf = prefix(add, x[0:n/2])
 secondhalf = prefix(add, x[n/2+1:n])
 secondhalf = map(add(x[n/2]), secondhalf)
 res = firsthalf + secondhalf
 return res
```

- Is this work efficient??

# Recursive Odd/Evens

- Suppose we extracted the odds and even indexed elements from an array and ran scan in parallel on both. Are we near done??
- Improved Complexity?
- ```
def extract_even_indexed (x):  
    result = range(len(x)/2)  
    indx = range(0,len(x)-1,2)  
    for i in indx:  
        result[i/2]=x[i]  
    return result
```


A Work-efficient Solution

```
def prefix(f,x):  
    if len(x) == 1:  
        return x  
    else: #begin parallel  
        e = extract_even_indexed(x)  
        o = extract_odd_indexed(x) #end parallel  
        s = map(f,e,o)  
        r1 = [0] + prefix(f,s)  
        r2 = map(f,e+[0],r1)  
        return interleave(r2[0:len(r2)-1],r1[1:])
```

Segmented Scan

- Indication of segments within array
- Apply scan operation to segments independently
- Work an example using both inclusive and exclusive scans.
- Next: application to sparse matrices

Sparse Matrix Dense Vector Products

- *Compressed Sparse Row* CSR format
- Value = nonzero data one array
- Column= identifies column for each
- Rowptr= pointers to location of each 1st data in each row

$$\begin{bmatrix} a & 0 & b \\ c & d & e \\ 0 & 0 & f \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix} =$$

VALUE $[a \ b \ c \ d \ e \ f]$
 COLUMN $[0 \ 2 \ 0 \ 1 \ 2 \ 1]$
 ROWPTR $[0 \ 2 \ 5]$

VECTOR
 $\begin{bmatrix} x \\ y \\ z \end{bmatrix}$

1. CREATE SEGMENTED REF'N FROM VALUE + ROWPTR
2. GATHER VECTOR VALUES USING COLUMN
3. PAIRWISE MULTIPLY 1.2

$\begin{bmatrix} a & b & | & c & d & e & | & f \end{bmatrix}$
 $\begin{bmatrix} x & z & & x & y & z & & y \end{bmatrix}$

1 CREATE SEGMENTED REF'N FROM
VALUE + ROWPTR

2 GATHER VECTOR VALUES USING
COLUMN

3 PAIRWISE MULTIPLY 1.2
(BACKWARDS)

4 EXCLUSIVE SEGMENTED SUM SCAN

$$\begin{bmatrix} a & b & | & c & d & e & | & f \end{bmatrix}$$

↑ ↑ ↑

$$\begin{bmatrix} x & z & & x & y & z & & y \end{bmatrix}$$

$$\begin{bmatrix} a \cdot x & b \cdot z & | & c \cdot x & d \cdot y & e \cdot z & | & f \cdot y \end{bmatrix}$$

+ + +

out(0) out(1) out(2)

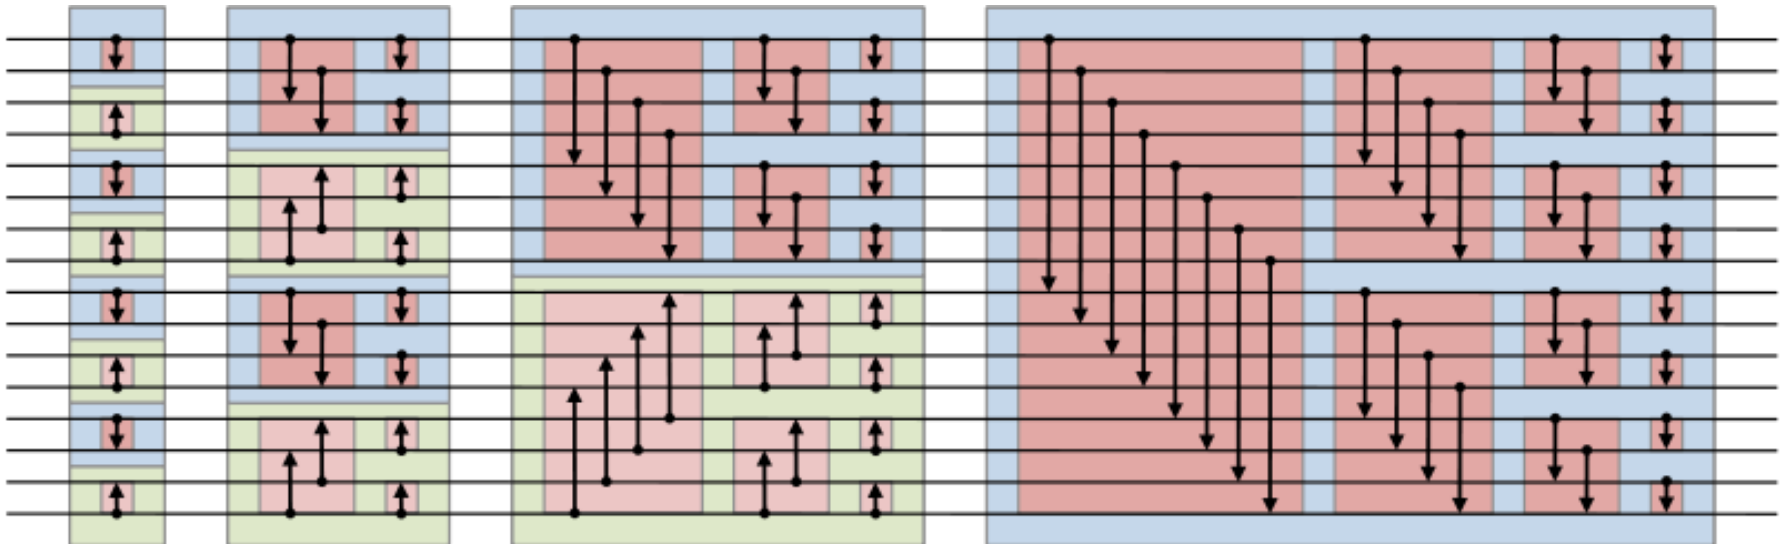
$$\begin{bmatrix} a & 0 & b \\ c & d & e \\ 0 & 0 & f \end{bmatrix} \begin{bmatrix} x \\ y \\ z \end{bmatrix} = \begin{bmatrix} ax + \cancel{0y} + bz \\ cx + dy + ez \\ \cancel{0x} + \cancel{0y} + fz \end{bmatrix}$$

Sorting

- Studied for classic patterns of data communication
- First look at favorite sequential algorithms and discuss parallelization
- BubbleSort
- QuickSort
- MergeSort
 - Three phase strategy

Sorting Networks

- Oblivious algorithms easily mapped to GPUs
- Bitonic Sorting
 - Ken Batcher (Kent State, Goodyear (Akron))
 - Bitonic Sequence (defined as a sequence of numbers with one direction change)



Recursive Bitonic Sorting

- ```
def bitonic_sort(up, x):
 if len(x) <= 1:
 return x
 else:
 first = bitonic_sort(True, x[:len(x) / 2])
 second = bitonic_sort(False, x[len(x) / 2:])
 return bitonic_merge(up, first + second)
```



# Bitonic Merge

- `def bitonic_merge(up, x):`  
    # assume input x is bitonic  
    # sorted list is returned  
    if `len(x) == 1`: return x  
    else:  
        `bitonic_compare(up, x)`  
        `first = bitonic_merge(up, x[:len(x) / 2])`  
        `second = bitonic_merge(up, x[len(x) / 2:])`  
        return `first + second`
- `def bitonic_compare(up, x):`  
    `dist = len(x) / 2`  
    for i in `range(dist)`:  
        if `(x[i] > x[i + dist]) == up`:  
            `x[i], x[i + dist] = x[i + dist], x[i]`

# Proof that Bitonic Sorting is correct

Assume given as input 0/1 sequence (applying Knuth's 0/1 Sorting Principle)

Assume that the length of the sequence is a power of 2

If the sequence is of length 1, do nothing

Otherwise, proceed as follows:

- Split the bitonic 0/1 sequence of length  $n$  into the first half and the second half i.e. 0000...01111...100000...0
- Perform  $n/2$  compare interchange operations in parallel of the form  $(i, i + n/2)$ ,  $0 \leq i < n/2$  (i.e., between corresponding items of the two halves)
- Claim: Either the first half is all 0's and the second half is bitonic, or the first half is bitonic and the second half is all 1's
- Therefore, it is sufficient to apply the same construction recursively on the two halves - Done!

# Radix Sort – High Performance Sort

- Relies on numerical representation
- Example: Binary numbers
- Start with LSB move to MSB
- Each stage split into numbers into 2 sets
- Work Complexity
- Step Complexity
- Optimizations

# HW#4 Remove Red Eye Effect

Stencil – normalized cross correlation scoring (done!)

**Focus of HW4:**

Sort all pixels using ncc scores

Map Operation: to remove red  
from highest scoring pixels  
(done!)

