

CS6068

Parallel Computing

Lect 4 – Sept 22, 2015

Topics:

Udacity Unit 5

Optimization - Achieving Terascale

CUDA Memory Restrictions

Using Shared Memory

Running Matrix Mult Example

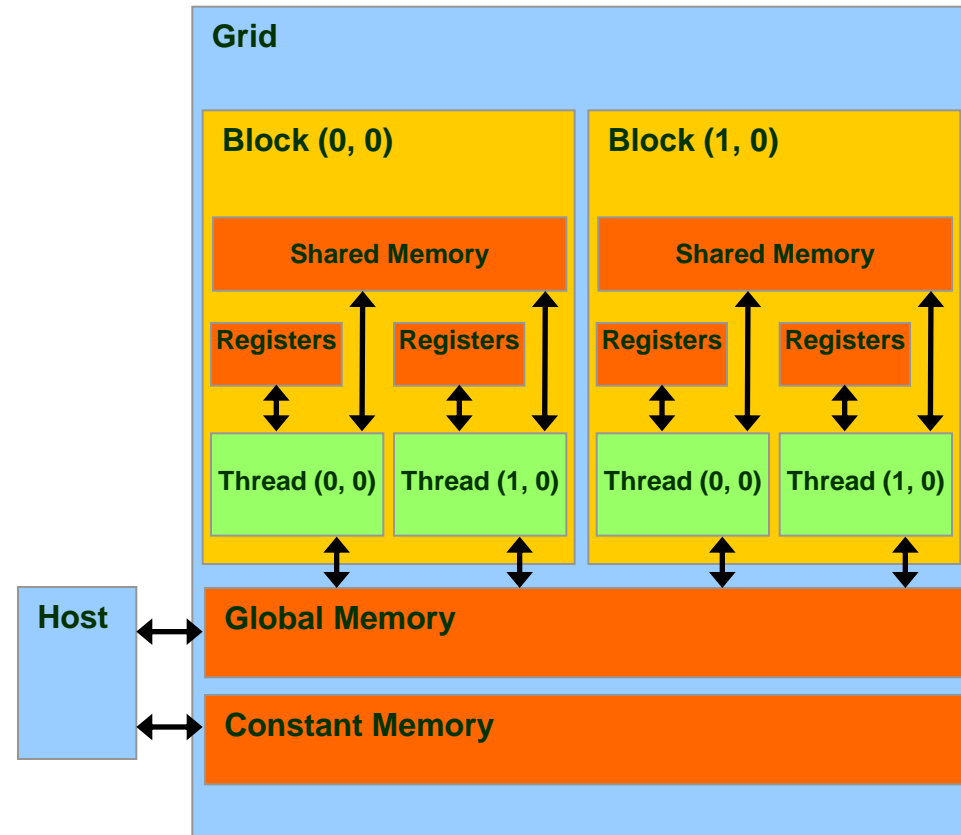
Performance Evaluation

Important Notions for Optimization

- First improve algorithm! REVIEW
- Second Look at Memory
- Most tuned GPU codes are Memory Limited.
- Device Query: Memory Clock Rate, Memory Bus Width → Theo Peak Bandwidth
- May use Profiler (like NVPP) to dig into potential Bottlenecks
- Memory Hierarchy, Coalescing, Ninja

Variety of CUDA Memories

- Each thread can:
 - Read/write per-thread **registers**
 - Read/write per-thread **local memory**
 - Read/write per-block **shared memory**
 - Read/write per-grid **global memory**
 - Read/only per-grid **constant memory**
 - Read/only per-grid **texture memory**



A Common Optimization: Tiling Data for Shared Memory

- Global memory resides in device memory (DRAM)
 - results in much slower access than shared memory
- **Tiling data** takes advantage of faster shared memory:
 - **Partition** data into **subsets** that fit into shared memory
 - Handle **each data subset with one thread block** by:
 - Load from global memory to shared memory,
 - Perform several computations
 - Copy results from shared memory to global memory

Another Common Optimization: Constant Memory

- Constant memory also resides in device DRAM - much slower than shared memory But...cached!
 - Highly efficient access for read-only data
- Carefully divide data according to access patterns
 - R/Only → constant memory (very fast if in cache)
 - R/W shared within Block → shared memory (very fast)
 - R/W within each thread → registers (very fast)
 - R/W inputs/results → global memory (very slow)

Optimization of Dot Products

Start design with relevant parameters

```
#define NumSMPs 32    #Machine dependent  
#define N 33 * 8192    #problem size  
#define ThreadsperBlock 256  
#define BlocksperGrid NumSMPs
```

```
__global__ void dot( int*a, int*b, int*c ) {  
    // Shared memory for results of multiplication  
  
    __shared__ int temp[N];  
    temp[threadIdx.x] = a[threadIdx.x]*b[threadIdx.x];  
  
    // Thread 0 sums the pairwise products  
  
    if( 0 == threadIdx.x ) {  
        intsum = 0;  
        for( int i= 0; i< N; i++ )  
            sum += temp[i];  
        *c = sum;  
    }  
}
```


Let's Deal with case where N is bigger than allowable number of threads ($\text{ThreadsperBlock} * \text{BlocksperGrid} = \text{blockDim.x} * \text{gridDim.x}$)

Our attempt begins with shared cache and initializing tid as offset into the N-item array.

```
__global__ void dot( float *a, float *b, float *c ) {  
    __shared__ float cache[ThreadsPerBlock];  
    int tid = threadIdx.x + blockIdx.x * blockDim.x;  
    int cacheIdx = threadIdx.x;  
    float temp = 0;  
    while (tid < N) {  
        temp += a[tid] * b[tid];  
        tid += blockDim.x * gridDim.x;  
    }  
    cache[cacheIdx] = temp; }  
}
```

Still need to reduce each block, and reduce across blocks

```
//add pairs of values offset by 2-powers  
// recursive halving
```

```
__shared__ float cache[ThreadsPerBlock];  
int cacheIdx = threadIdx.x;  
int offset = blockDim.x/2;  
while (offset != 0) {  
    if (cacheIdx < offset){  
        cache[cacheIdx] += cache[cacheIdx +  
offset];}  
    __syncthreads();  
    offset /= 2; }  
if (cacheIdx == 0)  
c[blockIdx.x] = cache[0];  
}
```

Matrix Multiplication Optimization

GOAL: ACHIEVING TERASCALE for Matrix Multiply

Assume that all threads access global memory for their input matrix elements

- Two memory accesses (8 bytes) per floating point multiply-add

- 4B/s of memory bandwidth/

FLOPS

- Tesla has 1030 GFLOP peak

- $4 \times 1030 = 4120$ GB/s required to achieve peak FLOP rating

- 148 GB/s on global memory limits the code at 37 GFLOPS

- Need to drastically cut down global memory accesses to get closer to the terascale peak 1030 GFLOPS

```
// Matrix multiply on host in double
precision
void MatrixMulOnHost(float* M, float* N,
float* P, int Width)
{
for (int i = 0; i < Width; ++i)
    for (int j = 0; j < Width; ++j) {
        double sum = 0;
        for (int k = 0; k < Width; ++k) {
            double a = M[i * width + k];
            double b = N[k * width + j];
            sum += a * b;
        }
        P[i * Width + j] = sum;
    }
}
```

```

void MatrixMulOnDevice(float* M, float* N, float* P, int
Width)
{
int size = Width * Width * sizeof(float);
float* Md, Nd, Pd;
// Allocate and Load M, N to device memory
cudaMalloc(&Md, size);
cudaMemcpy(Md, M, size, cudaMemcpyHostToDevice);
cudaMalloc(&Nd, size);
cudaMemcpy(Nd, N, size, cudaMemcpyHostToDevice);
// Allocate P on the device
cudaMalloc(&Pd, size);

// Kernel invocation code - to be shown later
...
// Read P from the device
cudaMemcpy(P, Pd, size, cudaMemcpyDeviceToHost);
// Free device matrices
cudaFree(Md); cudaFree(Nd); cudaFree (Pd);
}

```

MM: Kernel Invocation with 1 Block

```
// Setup the execution configuration
#define WIDTH 8192
dim3 dimGrid(1, 1);
dim3 dimBlock(WIDTH, WIDTH);
// Launch the device computation threads!
// 1x1 block with 8192x8192 threads
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd,
Pd,WIDTH);
--- MP Information for device 4 ---
Multiprocessor count: 30
Shared mem per mp: 16384
Registers per mp: 16384
Threads in warp: 32
Max threads per block: 512
Max thread dimensions: (512, 512, 64)
Max grid dimensions: (65535, 65535, 1)
```

MM: Kernel Invocation

```
// Setup the execution configuration
#define WIDTH 8192
#define TILE_WIDTH 16


dim3 dimGrid(WIDTH/TILE_WIDTH, WIDTH/TILE_WIDTH);
dim3 dimBlock(TILE_WIDTH, TILE_WIDTH);

// Launch the device computation threads!
// 512x512 blocks with 16x16 threads per block
MatrixMulKernel<<<dimGrid, dimBlock>>>(Md, Nd,
Pd,WIDTH);
```


Extend to Arbitrary Sized Square Matrices

- Have each 2D thread block to compute a $(\text{TILE_WIDTH})^2$ sub-matrix (tile) of the result matrix
 - Each has $(\text{TILE_WIDTH})^2$ threads
- Generate a 2D Grid of $(\text{WIDTH}/\text{TILE_WIDTH})^2$ blocks

You may still need to put a loop around the kernel call for cases where $\text{WIDTH}/\text{TILE_WIDTH}$ is greater than max grid size (64K)!



Matrix Multiplication Kernel using Multiple Blocks - Tiling

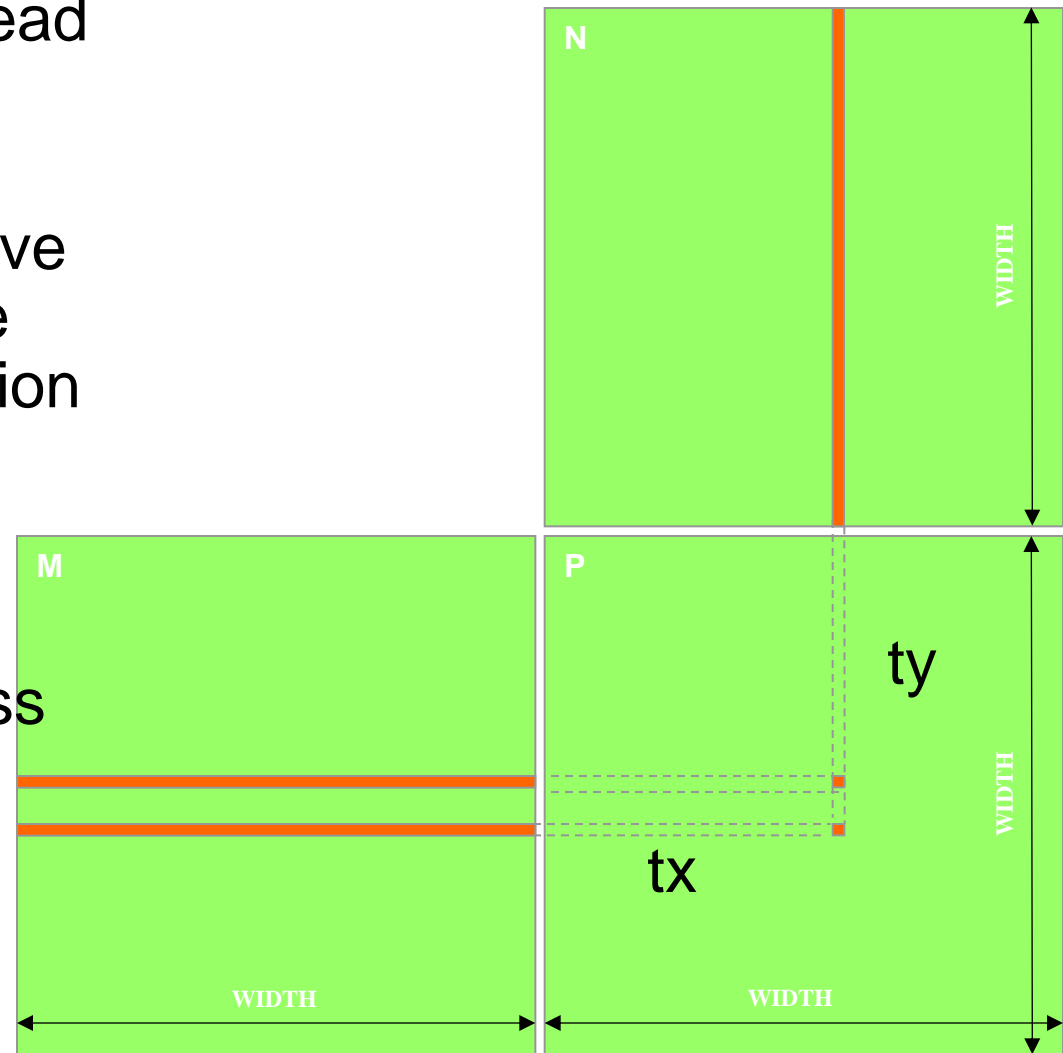
```
__global__ void MatrixMulKernel
    (float* Md, float* Nd, float* Pd, int Width)
{
    // Calculate the row index of the Pd element and M
    int Row = blockIdx.y*TILE_WIDTH + threadIdx.y;
    // Calculate the column index of Pd and N
    int Col = blockIdx.x*TILE_WIDTH + threadIdx.x;

    float Pvalue = 0;
    // each thread computes one element of the block sub-matrix
    for (int k = 0; k < Width; ++k)
        Pvalue += Md[Row*Width+k] * Nd[k*Width+Col];

    Pd[Row*Width+Col] = Pvalue;
}
```

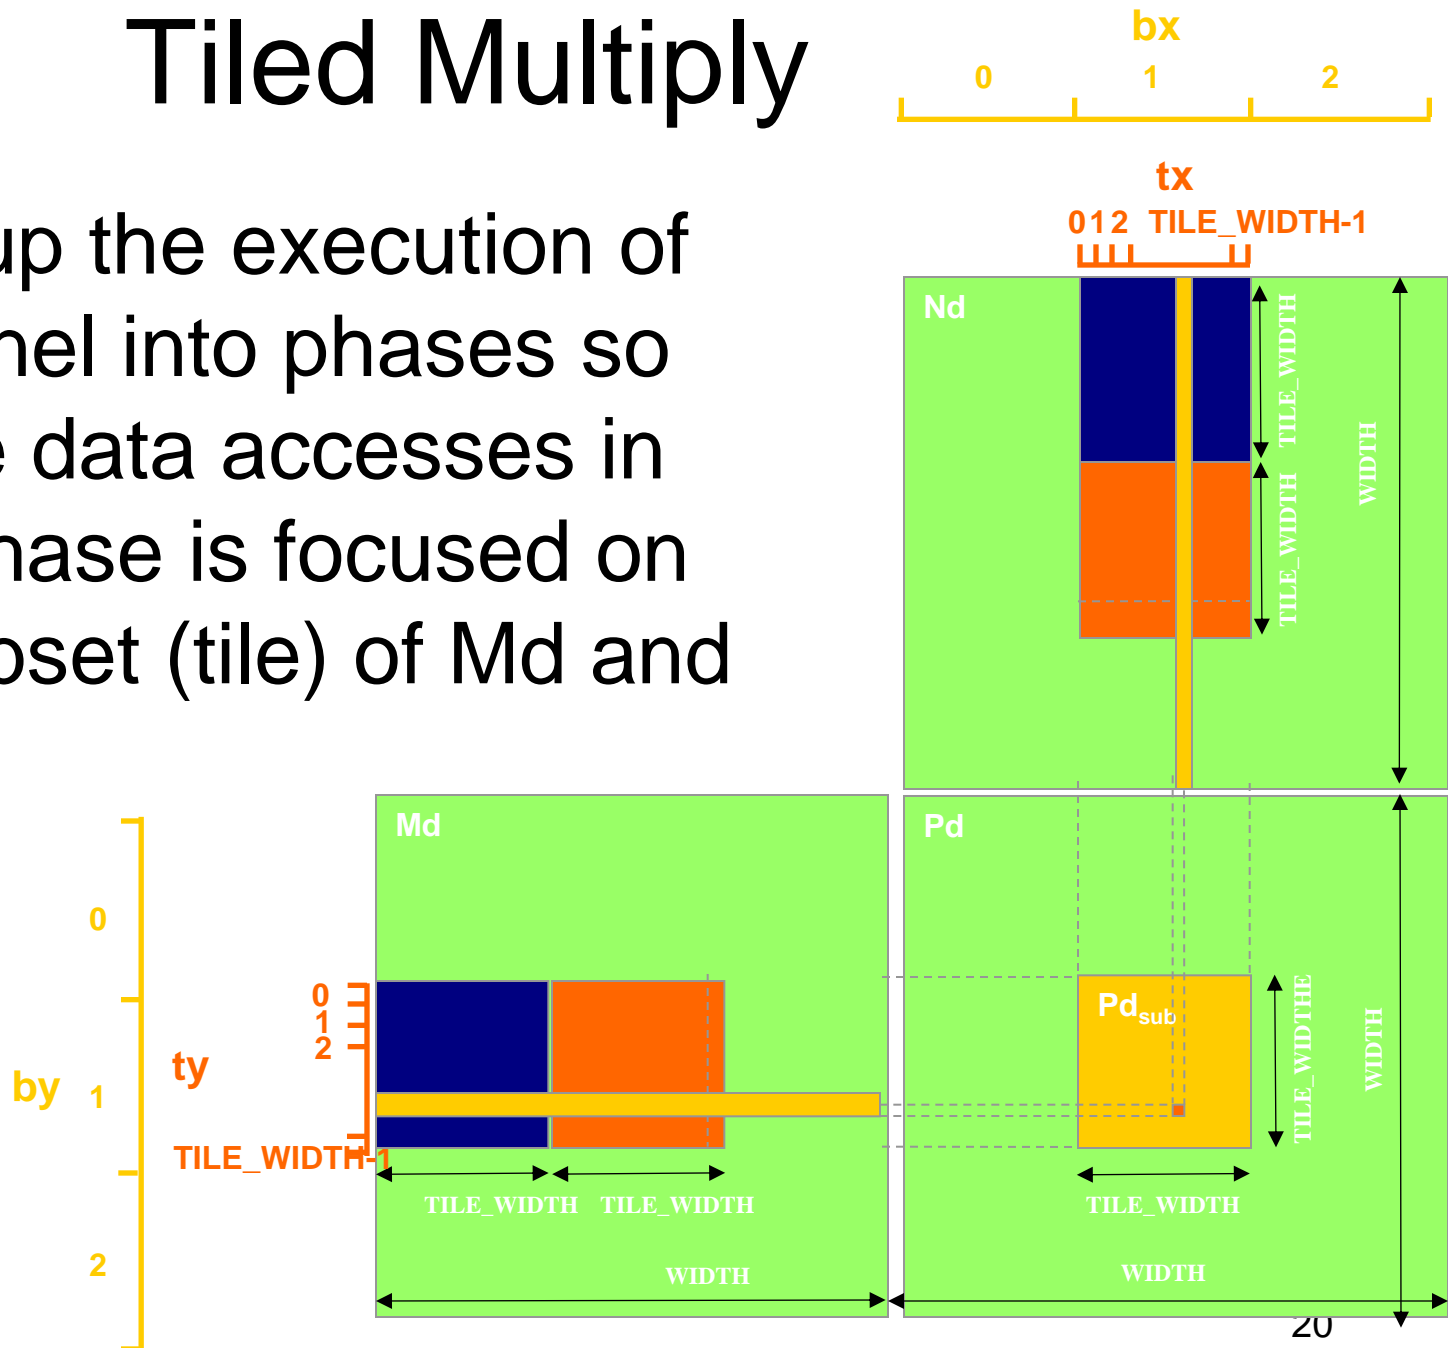
Idea: Reuse limited shared memory

- Each input element is read by Width threads.
- Load each element into Shared Memory and have several threads in same block use the local version to reduce the memory bandwidth
 - Tiled algorithms
- Try for one global access for each tile. All other access shared.

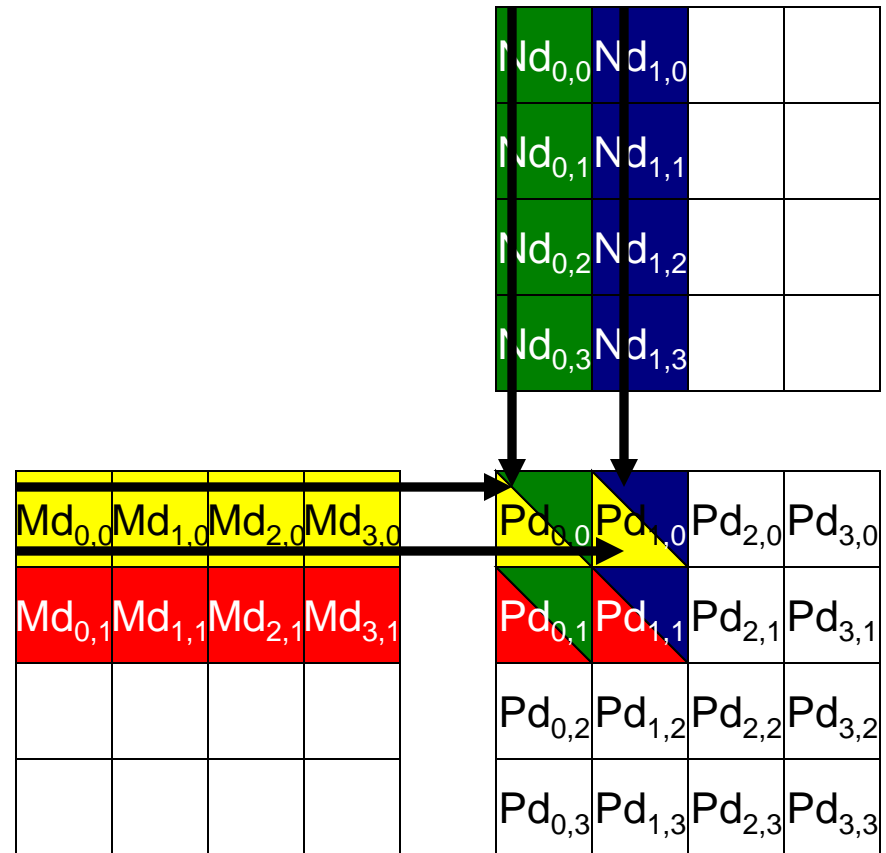


Tiled Multiply


- Break up the execution of the kernel into phases so that the data accesses in each phase is focused on one subset (tile) of M_d and N_d



Illustrate Idea with Small Example- 4 blocks of 4 threads each



Breaking Md and Nd into Tiles



Size Considerations in CUDA

- Each **thread block** should have many threads
 - TILE_WIDTH of 16 gives $16*16 = 256$ threads
- There should be many thread blocks
 - A $1024*1024$ Pd gives $64*64 = 4096$ Thread Blocks
- Each thread block perform $2*256 = 512$ float loads from global memory for $256 * (2*16) = 8,192$ mul/add operations.
 - 16 Flops per load
 - Global Memory bandwidth no longer a limiting factor

Reference Code:

Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* Md, float* Nd,
    float* Pd, int Width)
{
1.  __shared__ float Mds[TILE_WIDTH][TILE_WIDTH];
2.  __shared__ float Nds[TILE_WIDTH][TILE_WIDTH];

3.  int bx = blockIdx.x;  int by = blockIdx.y;
4.  int tx = threadIdx.x; int ty = threadIdx.y;

// Identify the row and column of the Pd element to
  work on
5.  int Row = by * TILE_WIDTH + ty;
6.  int Col = bx * TILE_WIDTH + tx;
```


Reference Code:

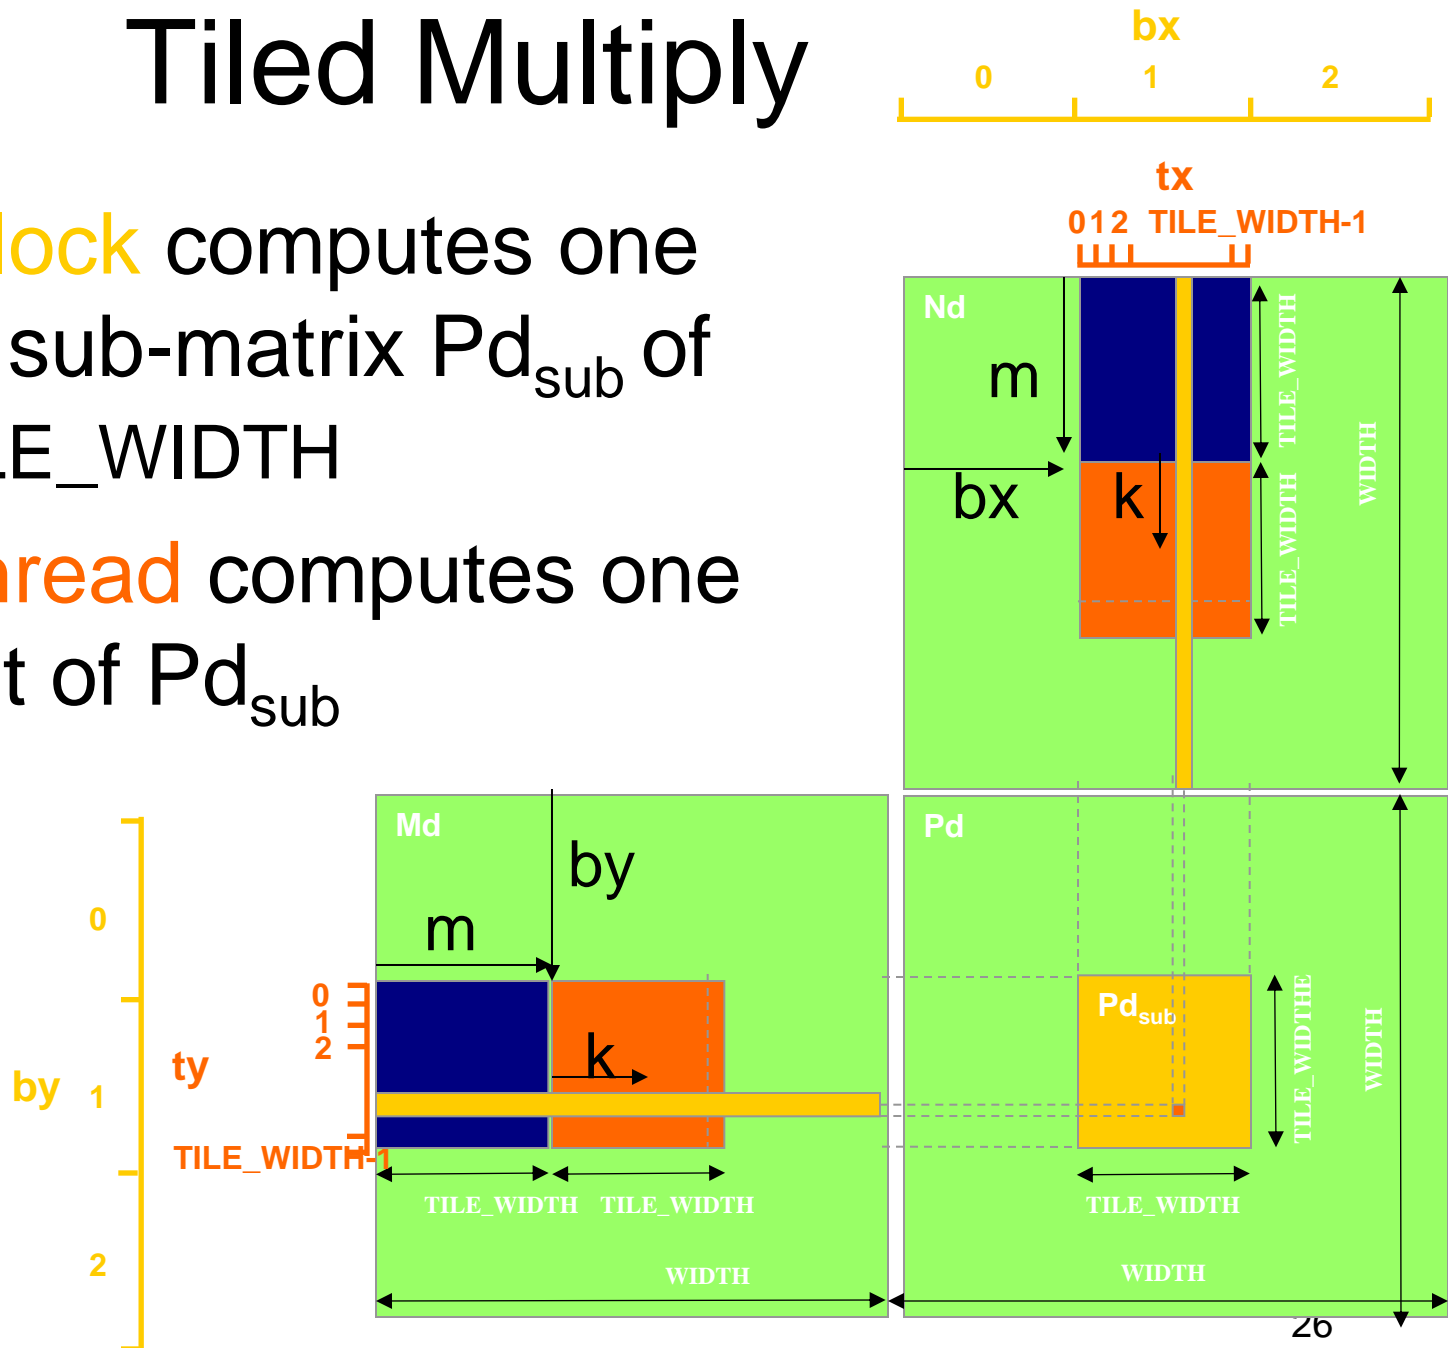
Tiled Matrix Multiplication Kernel

```
float Pvalue = 0;
// Loop over the Md and Nd tiles required to compute
// the Pd element
7.   for (int m = 0; m < Width/TILE_WIDTH; ++m) {
// Collaborative loading of Md and Nd tiles into shared
// memory
8.   Mds[ty][tx] = Md[Row*Width + (m*TILE_WIDTH + tx)];
    Nds[ty][tx] = Nd[Col + (m*TILE_WIDTH + ty)*Width];
    __syncthreads();

9.   for (int k = 0; k < TILE_WIDTH; ++k)
        Pvalue += Mds[ty][k] * Nds[k][tx];
        __syncthreads();
}
10.  Pd[Row*Width+Col] = Pvalue;
}
```

Tiled Multiply

- Each **block** computes one square sub-matrix Pd_{sub} of size $TILE_WIDTH$
- Each **thread** computes one element of Pd_{sub}



CUDA Shared Memory and Latency Hiding

- Each SM has at least 16KB shared memory
 - SM size is implementation dependent!
 - For $\text{TILE_WIDTH} = 16$, each thread block uses $2 \times 256 \times 4\text{B} = 2\text{KB}$ of shared memory.
 - Can potentially have at least 8 Thread Blocks actively executing
 - This allows up to $8 \times 512 = 4,096$ pending loads. (2 per thread, 256 threads per block)
 - The next TILE_WIDTH 32 would lead to $2 \times 32 \times 32 \times 4\text{B} = 8\text{KB}$ shared memory usage per thread block, allowing only up to two thread blocks active at the same time
- Using 16x16 tiling, we reduce the accesses to the global memory by a factor of 16
 - The 148 B/s bandwidth can now support $(148/4) \times 16 = 592$ GFLOPS!

Ninja Optimizations:


- WARP Control Flow and Divergence
- Bank Conflicts

Important:

CUDA runs smoother when consecutively ordered threads are doing same operation

- Consider the problem of getting all threads in a block to the same state, such as summing all items in cache array. Doing a reduction followed by broadcast creates divergence.

- Each level of the butterfly links indices that differ in single *bit* position



```
__shared__ int bfly[ThreadsperBlock];
int cacheIdx = threadIdx.x;

bfly[cacheIdx] = cache[cacheIdx] ;
__syncthreads();
for(int bit=blocksize/2; bit>0; bit/=2)
{
    int tmp = bfly[cacheIdx]+ bfly[cacheIdx^bit];
    __syncthreads();
    bfly[cacheIdx]=tmp;
    __syncthreads();
}
// Every bfly item now holds sum of all cache
```