

# CS6068

Parallel Computing - Fall 2014

Lecture Week 6 – Oct 13

Topics:

Unit 6 Udacity Continues

Review Breadth-First Graph Traversals

Depth-first and P-complete problems

Processing Linked Lists in Parallel

Trading Work for Steps

Bloom Filters and Cuckoo Hashing

# Parallel Graph Traversal

WWW, Facebook, Tor

Application: Visit every node once

Breadth-First Traversal:

Visit nodes level-by-level, synchronously

Depth-First Traversal:

Visit nodes at periphery first.

Understand Variety of Graphs : Small Depth,  
Large Depth, Small-world

# Design of BFS

Goal: Compute hop distance of every node from source node.

1<sup>st</sup> try:

Thread per Edge Method

Work Complexity

Step Complexity

Control Iterations


Race Conditions

Finishing Conditions

INITIALIZE: STARTING NODE'S DEPTH = 0


STARTING FRONTIER = NEIGHBORS OF  
STARTING NODE

- 1 FRONTIER: FIND NEIGHBOR START
- 2 : HOW MANY NEIGHBORS?
- 3 ALLOCATE SPACE FOR NEW FRONTIER
- 4 COPY EDGE LIST TO NEW ARRAY
- 5 CULL VISITED ELEMENTS




# Parallel BFS Strategies

1. Expand current frontier (**level-synchronous** approach, suited for **low diameter** graphs)



- $O(D)$  parallel steps
- Adjacencies of all vertices in current frontier are visited in parallel

2. Stitch multiple concurrent traversals (Ullman-Yannakakis approach, suited for **high-diameter** graphs)



- path-limited searches from “super vertices”
- APSP between “super vertices”

# Depth-first Traversal

```
def depthFirst (v):  
    if marked[v]: return  
    else:  
        marked[v]=true  
        for each w in Neighbor(v):  
            depthFirst(w)
```

Depth-first traversal is used in many apps

- Topological sort of DAGs
- Detecting cycles in graphs
- Strongly-connected components

Depth-first ordering is labeling of vertices that are consistent with the ordering of a depth-first traversal.

The problem of computing depth-first ordering is P-complete.

# Parallel Complexity Theory

What problems are inherently sequential?

Recall the sequential complexity classes:

P, NP, NP-complete 1MQ:  $P = NP$  ?

A class to describe parallel problems: NC

- solvable in polylog time on poly number of processors

Now we can ask what are the hardest polytime sequential problems? Call this **P-complete**

Examples:

Circuit Value Problem

Conway's Game of Life

Lexicographically First Depth First Ordering

# Processing a Linked List

## Design Strategy: Recursive Pointer Jumping

The idea is to do pointer indirection in parallel and turn a single linked list into 2 shorter (half the length) linked lists in one operation, then recurse.

Here is the basic code:

```
def pointer_jump_rec (list-of-ptrs):  
    chum-ptrs = range(len(list-of-ptrs))  
    for i in range(len(chum-ptrs)):  
        chumptr[i] = list-of-ptrs[list-of-ptrs[i]]  
    if eql(chum-ptrs, list-of-ptrs):  
        return chum-ptrs  
    else:  
        return pointer_jump_rec(chum-ptrs)
```



## **An Example of using Pointer Jumping**

0-> 2->1-> 5-> 6      4-> 7-> 3

we have two linked lists that can be stored as one  
python list of index pointers

```
ptrs=[2, 5, 1, 3, 7, 6, 6, 3].
```

If we apply pointer jumping each node index will  
eventually point to the head of the list it is a member  
of.

```
>>> pointer_jump_rec ([2, 5, 1, 3, 7, 6, 6, 3])
```

```
[6, 6, 6, 3, 3, 6, 6, 3]
```

Work and Step complexity of Pointer-Jumping??

## **List Ranking Problem:**

Determine the hop distance of each node to its head node.

Can we apply Pointer-Jumping paradigm to solve List-Ranking Problem?

We can update the rank of node with each call to pointer-jump:

```
def jump_rank (ptrs, oldranks):
    newranks= range(len(oldranks))
    for i in range(len(ranks)):
        newranks[i] = oldranks[i] + oldranks[ptrs[i]]
    chum-ptrs = range(len(ptrs))
    for i in range(len(ptrs)):
        chum-ptrs[i] = ptrs[ptrs[i]]
    if eql(chumptrs,ptrs):
        return newranks
    else:
        return jump_rank(chum-ptrs,newranks)
```

Why does this work? What are initial ranks?

What is the invariant property of (chum-ptrs,newranks)

# Trading-off Work for Steps

List ranking is a classic example that shows that we can trade-off increasing work for reduction in step complexity.

Let's apply this principle to sorting. Suppose we permit unlimited work. How much can we reduce the step complexity of sorting??

# Counting Sort

```
def countingSort(list):  
    scatter = range(len(list))  
    result = range(len(list))  
    for i in range(len(list)): #do for loop in parallel  
        scatter[i] = len(compact(lessthan(list[i]), list))  
        result[scatter[i]] = list[i]  
    return result
```

What is the work and step complexity of countingSort??

# Parallel Hashing

Hashing Problem: Fast lookup in table using keys.

Applications: Implementing Disjoint Sets, Graph lookups

Static Case: Perfect Hashing

Dynamic: Insertions and deletions, bucket contention and collisions

Load Factors – if your hashing  $n$  items to  $m$  buckets what is optimal load factor?

Separate Chaining

Open Addressing

Problems with Chaining and Open Addressing in parallel.

# Bloom Filters for Set Membership

Problem: Fast test of inclusion – is given element already in a given set.

Bloom Filter BF Solution:

Select k-hash functions map to locations in BF table – large table of bits – initially all 0

Insert item into set by setting k-bits in BF using k-hashes

Testing for set inclusion checks k-bits

Small probability of false positives.

# Cuckoo Hashing – avoid problems with chaining

Select  $k > 1$  hash functions

To insert item - Map item to any of the  $k$  locations

If there are  $k$ -collisions we bump one from nest.

Reinsert the bumped element.

When can this fail? Under what conditions?

Is it likely to fail? How large do we choose  $n$ ?



k	h(k)	h'(k)
20	9	1
50	6	4
53	9	4
75	9	6
100	1	9
67	1	6
105	6	9
3	3	0
36	3	3
39	6	3

1. table for h(k)											
	20	50	53	75	100	67	105	3	36	39	
0											
1					100	67	67	67	67	100	
2											
3								3	3	36	
4											
5											
6		50	50	50	50	50	105	105	105	50	
7											
8											
9	20	20	20	20	20	20	53	53	53	75	
10											

2. table for h'(k)											
	20	50	53	75	100	67	105	3	36	39	
0										3	
1							20	20	20	20	
2											
3									36	39	
4			53	53	53	53	50	50	50	53	
5											
6				75	75	75	75	75	75	67	
7											
8											
9						100	100	100	100	105	
10											